

AUSHANG · THE PAPER NOTICE BOARD, DIGITIZED

By every Kita door hangs a board no parent has time to read.

Aushang turns one photo of that board into a feed, a calendar and reminders on every parent's phone — in German, English or Russian. Staff change nothing: they keep the board, snap a picture, and the rest is automatic.

THE PROBLEM

Three traps of the paper board

01

Information is trapped

Miss the board, miss the closure, the Thursday trip, the lice warning. No feed, no calendar, no digest — the information dies in a hallway.

02

Every app is a second system

The off-the-shelf answer makes staff retype every notice into a new tool. One overworked admin won't maintain it — so they stay on paper.

03

The data is sensitive

Notices name children, parents, phone numbers. Pushing a board photo into a generic cloud tool is a privacy problem these orgs are rightly nervous about.

Aushang's answer

It changes **none** of this process — it enhances it. Staff keep their paper board; one photo feeds the AI and the automation does the rest. No retyping, no second system.

Two principles, enforced by architecture

01 Privacy by construction

Every child's name, parent, phone number and address is detected and **masked locally** — on our own infrastructure — before any external AI call is made. Redaction is **fail-closed** : when the system is unsure, it masks.

Raw photographs never leave the organization's infrastructure, and no un-redacted personal data ever reaches the AI. This is enforced at the architecture layer — not by a written promise.

02 The AI advises, a human decides

The AI only **suggests** — a content type and the fields it can read from the redacted text. Every suggestion is checked against a strict schema; anything malformed routes to a manual path and **never auto-publishes**.

An admin reviews and confirms before a post reaches a single member — and the system routes only on the **confirmed** value, never on the AI's raw guess.

Together, these two principles are what keep Aushang out of EU AI Act **High-risk**.

THE PIPELINE

From photograph to published post



The only text that crosses to the AI is **redacted**. The only path to a member's screen is an **admin publish**.

IN PLAIN WORDS

OCR	Tesseract	Local redact	Presidio · spaCy	fail-closed	Fan-out	ICS
Reads the board's text from the photo.	The free engine doing that (German).	Blanks personal data on our own server.	Find names, parents & numbers to blank.	If unsure, it hides — never exposes.	One post sent to every channel at once.	Calendar subscription; dates land in the phone.

PROOF OF CONCEPT

Proven before building the real app

01 · Personal data

Personal data is removed before the AI sees anything

Our own code masks the names and numbers; Claude only ever receives the cleaned text.

02 · No bad answers

A bad AI answer can never go live

If the AI returns something malformed, it's sent to a person — it is never published automatically.

03 · Human decides

The human's choice decides, not the AI's guess

Where a post goes follows what the admin confirmed — change the type, the destination changes.


04 · Photo consent

A photo needs two yeses to appear

A clear photo is shown only if a parent opts in **and** an admin releases it.

THE LIVE SYSTEM · NOT A MOCKUP


A real Kita, today



Capture


Photograph the board, or upload.

kita-connect.cloud · invite-only



Private feed


Typed posts — Termin, Info, Rückblick.



Calendar + ICS

Subscribe once; dates sync to the phone.

German UI today · English & Russian in product · Android shell builds in CI



Members

Invite-only, groups, QR onboarding.


Contribution-positive from a low base

PER YEAR · €	5 Kitas	10 Kitas	20 Kitas · base	50 Kitas
Subscriptions · €15/mo	900	1.800	3.600	9.000
App sales · net of fees	383	765	1.530	3.825
Revenue	1.283	2.565	5.130	12.825
Infrastructure +LLM	-1.360	-1.360	-1.360	-1.360
Store licences · yr 1	-115	-115	-115	-115
Net contribution	-192	+1.090	+3.655	+11.350

Assumptions — €15/Kita/mo; ~1 paid app download per Kita per month, blended €6,38 net of the 15% store fee (Apple €9→€7,65 · Google €6→€5,10); infrastructure ≈flat at this scale; licences Apple €90/yr + Google €25 once; build cost (sweat equity) excluded. Cash break-even ≈8 Kitas.

RISK · LIKELIHOOD × IMPACT

Eight risks, scored and mitigated



R4 · Low adoption / distribution

16

No-process-change design + a channel (association / Träger) is the make-or-break.

R1 · PII leak to the AI

15

Local fail-closed redaction + DB column REVOKE + double-gated consent — the core thesis.

R2 · LLM hallucination

12

Schema-validate + manual fallback + admin confirms every publish.



R3 over-redaction · R5 infra failure · R6 surveillance creep · R7 AI-Act misclassification · R8 LLM residency — all scored ≤9, each degrades to a safe manual path.

PII =personal data — names, parents' phone numbers, children's birthdates.

EU AI Act + GDPR

EU AI Act LIMITED RISK Makes no decision about a person — it classifies a notice, not a learner. No Annex III area. ✓ 100% of posts need explicit admin confirmation (human oversight). ✓ AI suggestions are labelled before confirm (Art. 50 transparency). ✓ Schema-validated, fail-closed, with a manual path (robustness).	GDPR RESIDUAL LOW-MED Only locally-redacted text reaches the AI — no PII, no raw image. ✓ Raw photos stay in EU infrastructure (Supabase EU). ✓ DPIA completed on the highest-risk activity. ✓ Data-subject rights: access, erasure, rectification, portability.
---	---

Roadmap and go-to-market



NOW · BOTTOM-UP

Start where it's real

The pilot runs in my own child's Kita. Every additional Kita pays **€15/mo just to cover infra** — break-even around 15 orgs. Cheap, honest, and it builds the reference.

SCALE · CHANNEL

Then sell through associations

A Kita-association or Träger represents dozens of orgs. One sale onboards many at €15–25/org/mo — this is what repays the build. The DPO-ready privacy story makes the sale.

WHY WE WIN

- No process change — just photograph the board.
- Privacy by construction — the DPO says yes.
- A live, compliant reference build — not a mockup.

AI advises. Personal data is masked first. A human decides.

Proven on a live, GDPR-by-construction build.

• **Verdict — ship & validate, then pursue the channel**

Q & A

Does a child's name reach the AI?

No — masked locally, fail-closed, before the call.

What if the AI mis-reads?

Schema-validate routes to manual; admin confirms.

Will small Kitas adopt it?

No process change — just photograph the board.

Why Limited risk, not High?

No Annex III area — no decision about a person.

US-hosted LLM — GDPR issue?

Only redacted text crosses; EU swap is one module.

Is it really live?

Yes — kita-connect.cloud, first Kita testing now.

AI advises. Personal data is masked first. A human decides.

Proven on a live, GDPR-by-construction build.

● **Verdict — ship & validate, then pursue the channel**

Q & A

Does a child's name reach the AI?

No — masked locally, fail-closed, before the call.

What if the AI mis-reads?

Schema-validate routes to manual; admin confirms.

Will small Kitas adopt it?

No process change — just photograph the board.

Why Limited risk, not High?

No Annex III area — no decision about a person.

US-hosted LLM — GDPR issue?

Only redacted text crosses; EU swap is one module.

Is it really live?

Yes — kita-connect.cloud, first Kita testing now.