

DIGITIZING THE PAPER NOTICE BOARD · WITHOUT CHANGING THE PROCESS

An AI reads the board. Personal data is masked first. A human decides.

It reads the photographed board, masks every name, parent and phone number **locally**, and suggests a structured post — in German, English or Russian. An admin confirms before anything publishes, and a super-admin can remove any post at any time.

THE PROBLEM

Three traps of the paper board

01

Information is trapped

Miss the board, miss the closure, the Thursday trip, the lice warning. No feed, no calendar, no digest — the information dies in a hallway.

02

Every app is a second system

The off-the-shelf answer makes staff retype every notice into a new tool. One overworked admin won't maintain it — so they stay on paper.

03

The data is sensitive

Notices name children, parents, phone numbers. Pushing a board photo into a generic cloud tool is a privacy problem these orgs are rightly nervous about.

Aushang's answer

It changes **none** of this process — it enhances it. Staff keep their paper board; one photo feeds the AI and the automation does the rest. No retyping, no second system.

Two principles, enforced by architecture

01 Privacy by construction

Every child's name, parent, phone number and address is detected and **masked locally** — on our own infrastructure — before any external AI call is made. Redaction is **fail-closed** : when the system is unsure, it masks.

Raw photographs never leave the organization's infrastructure, and no un-redacted personal data ever reaches the AI. This is enforced at the architecture layer — not by a written promise.

02 The AI advises, a human decides

The AI only **suggests** — a content type and the fields it can read from the redacted text. Every suggestion is checked against a strict schema; anything malformed routes to a manual path and **never auto-publishes**.

An admin reviews and confirms before a post reaches a single member — and the system routes only on the **confirmed** value, never on the AI's raw guess.

Together, these two principles are what keep Aushang out of EU AI Act **High-risk**.

THE PIPELINE

From photograph to published post



The only text that crosses to the AI is **redacted**. The only path to a member's screen is an **admin publish**.

ICS — the standard calendar feed (.ics). Members subscribe once; closures and dates sync to their own phone calendar, no app install needed.

Proven before a line of app code

01 · Redaction boundary

Proves a Code node owns masking — Claude only ever sees redacted text.

02 · Extract & validate

Proves malformed AI output routes to manual — it never auto-publishes.

03 · Confirm, route, publish

Proves routing reads the admin's confirmed value — flip the type, the route changes.

04 · Double-gated consent

Proves a clear photo releases only on member opt-in **and** admin release.

52

assertions · all green

0

lines of application code

Every workflow imports straight into n8n and runs on mock data.

Six mechanisms, not one promise

1 **Local fail-closed redaction**

Presidio + spaCy + German regex before any external call. On uncertainty, it masks.

2 **Database column REVOKE**

Raw-PII columns are revoked at the DB layer — even an admin's browser cannot read them.

3 **Double-gated photo consent**

A clear photo is released only on member opt-in and admin release, via short-TTL signed URLs.

4 **Reflection photo deleted at publish**

Rückblick originals — most likely to show children — are deleted from storage at publish.

5 **Build-blocking secret scan**

A leaked service-role or LLM key blocks the build before deploy. No key ships to the browser.


6 **Adversarial security review**

A four-layer model — middleware, helpers, security-definer RPCs, RLS — reviewed adversarially.

No raw image and no un-redacted PII is **ever** sent to the AI — redaction sits upstream of every external call, by architecture.


THE LIVE SYSTEM · NOT A MOCKUP

A real Kita, today




Capture

Photograph the board, or upload.




Private feed

Typed posts — Termin, Info, Rückblick.



Calendar + ICS

Subscribe once; dates sync to the phone.



Members

Invite-only, groups, QR onboarding.

kita-connect.cloud · invite-only

German UI today · English & Russian in product · Android shell builds in CI

Pricing, break-even, and the honest read

€15_{/mo}

per org — covers infra, nothing more

~15_{orgs}

cash break-even on infra

+177%

12-mo cash ROI · 20 orgs

−80%

36-mo all-in · incl. €35k build

How the money works

Cash infra / year	≈ €1.300
LLM (Claude Haiku, once per capture)	≈ €60
20 orgs × €15/mo revenue	€3.600
Net (12-mo, cash)	+€2.300

The first Kita — my own child's — runs the pilot at cost. Every additional Kita pays €15/mo just to cover infrastructure.

App-store line


A one-time €6 app purchase on Apple & Google — incremental revenue and a credibility signal for the channel.

The honest read

Cheap to run, feature proven. The lever is a **channel** , not technology: €25/org × 150 Kitas via an association ≈€45k/yr — the build repays inside year 1.

RISK · LIKELIHOOD × IMPACT

Eight risks, scored and mitigated



R4 · Low adoption / distribution

No-process-change design + a channel (association / Träger) is the make-or-break.

16

R1 · PII leak to the AI

Local fail-closed redaction + DB column REVOKE + double-gated consent — the core thesis.

15

R2 · LLM hallucination

Schema-validate + manual fallback + admin confirms every publish.



12

R3 over-redaction · R5 infra failure · R6 surveillance creep · R7 AI-Act misclassification · R8 LLM residency — all scored ≤9, each degrades to a safe manual path.

EU AI Act + GDPR

EU AI Act LIMITED RISK Makes no decision about a person — it classifies a notice, not a learner. No Annex III area. ✓ 100% of posts need explicit admin confirmation (human oversight). ✓ AI suggestions are labelled before confirm (Art. 50 transparency). ✓ Schema-validated, fail-closed, with a manual path (robustness).	GDPR RESIDUAL LOW-MED Only locally-redacted text reaches the AI — no PII, no raw image. ✓ Raw photos stay in EU infrastructure (Supabase EU). ✓ DPIA completed on the highest-risk activity. ✓ Data-subject rights: access, erasure, rectification, portability.
---	---

Roadmap and go-to-market



NOW · BOTTOM-UP

Start where it's real

The pilot runs in my own child's Kita. Every additional Kita pays **€15/mo just to cover infra** — break-even around 15 orgs. Cheap, honest, and it builds the reference.

SCALE · CHANNEL

Then sell through associations

A Kita-association or Träger represents dozens of orgs. One sale onboards many at €15–25/org/mo — this is what repays the build. The DPO-ready privacy story makes the sale.

WHY WE WIN

- No process change — just photograph the board.
- Privacy by construction — the DPO says yes.
- A live, compliant reference build — not a mockup.

AI advises. PII is masked first. A human decides.

Proven on a live, GDPR-by-construction build.

● **Verdict — ship & validate, then pursue the channel**

Q & A

Does a child's name reach the AI?

No — masked locally, fail-closed, before the call.

What if the AI mis-reads?

Schema-validate routes to manual; admin confirms.

Will small Kitas adopt it?

No process change — just photograph the board.

Why Limited risk, not High?

No Annex III area — no decision about a person.

US-hosted LLM — GDPR issue?

Only redacted text crosses; EU swap is one module.

Is it really live?

Yes — kita-connect.cloud, first Kita testing now.